



2021.09.23

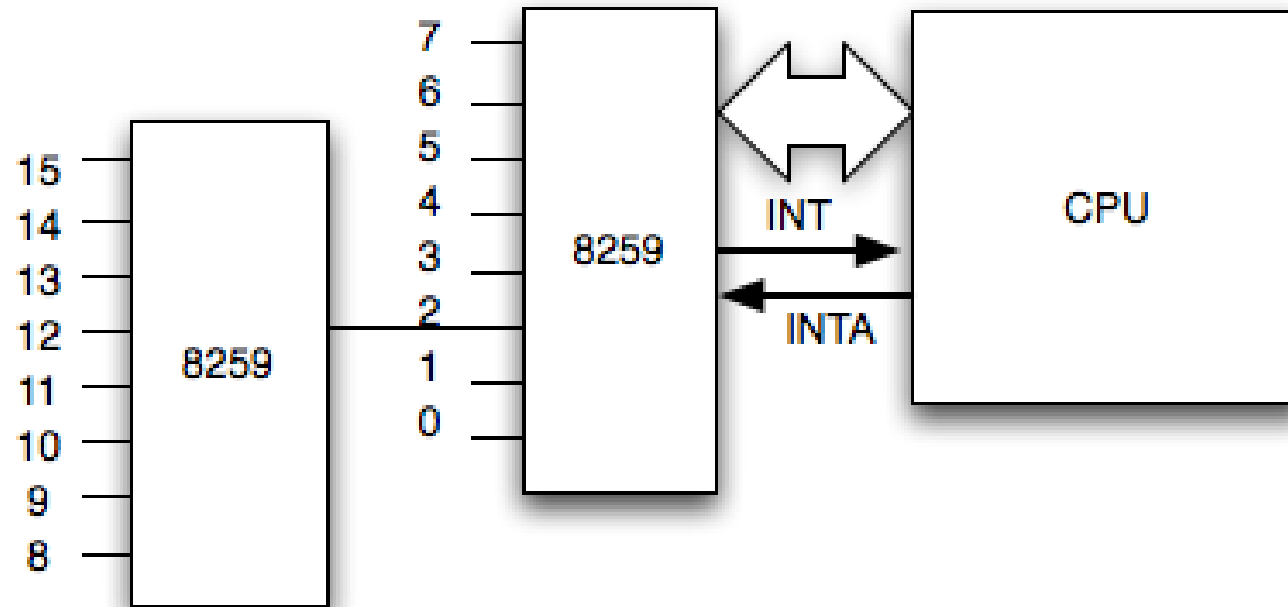


■ 彭天祥

Interrupt Controller

Interrupt Controller

Uniprocessor: Intel 8259A Interrupt Controller



Interrupt Controller

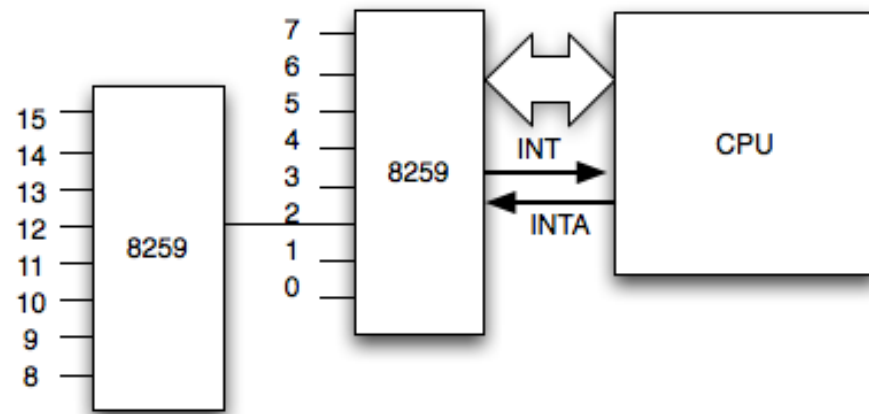
Uniprocessor: Intel 8259A Interrupt Controller

CPU side:

- INT向CPU assert中断
- 通过*ccli*屏蔽所有中断信号

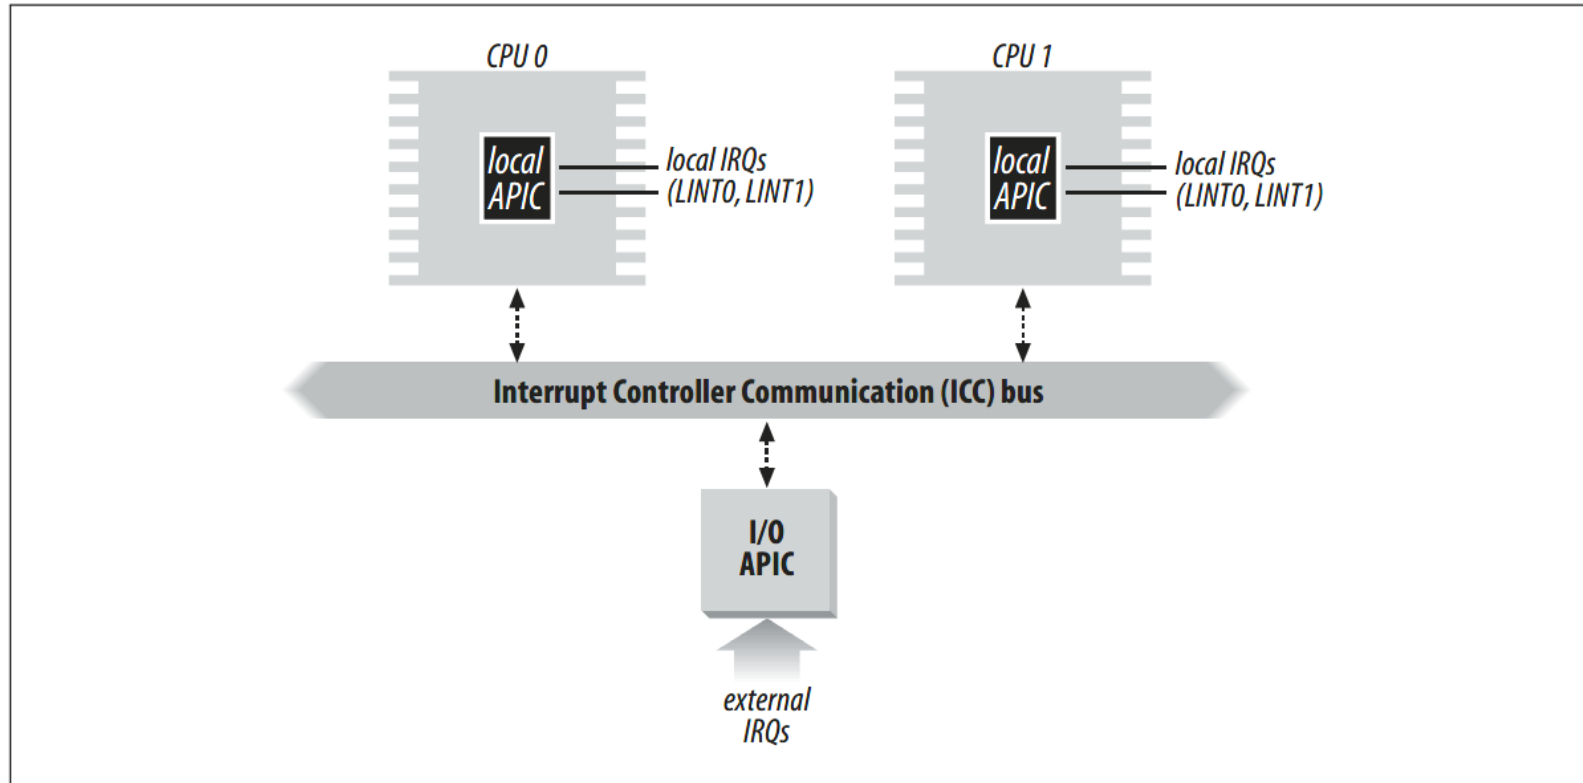
Device side:

- 多个IRQ pin
- 接收设备IRQ, 映射为Linux中的IRQ(+32), 向CPU side发送中断信号
- mask特定的IRQ line来仅屏蔽一条IRQ line上的所有中断



Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)



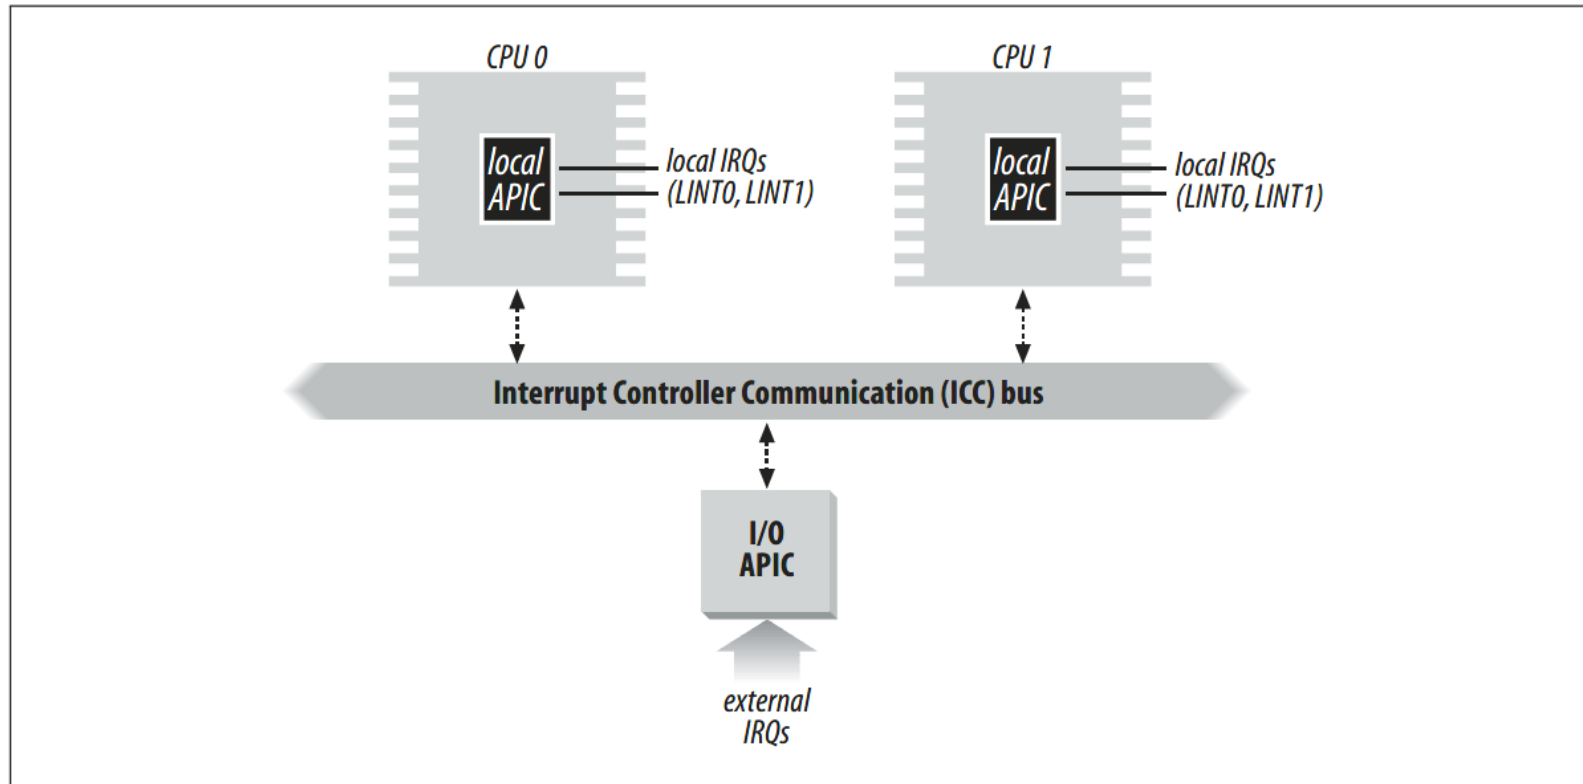
Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）
- 转发策略：
 - 静态：根据Redirection Table，可以一个/一些/广播

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）
- 转发策略：
 - 静态：根据Redirection Table，可以一个/一些/广播
 - 动态：每个Local APIC有一个硬件task priority register，希望os每次reschedule的时候更新，这样每个dynamic IRQ到来时分发到最低优先级的CPU，优先级相同时RR

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）
- 转发策略：
 - 静态：根据Redirection Table，可以一个/一些/广播
 - 动态：每个Local APIC有一个硬件task priority register，希望os每次reschedule的时候更新，这样每个dynamic IRQ到来时分发到最低优先级的CPU，优先级相同时RR
- cli时仅屏蔽local APIC中断

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)




Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）
- 转发策略：
 - 静态：根据Redirection Table，可以一个/一些/广播
 - 动态：每个Local APIC有一个硬件task priority register，希望os每次reschedule的时候更新，这样每个dynamic IRQ到来时分发到最低优先级的CPU，优先级相同时RR
- cli时仅屏蔽local APIC中断
- **mask一条IRQ line时会使I/O APIC不再响应这个irq，从而对所有CPU屏蔽掉一条IRQ**

Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)



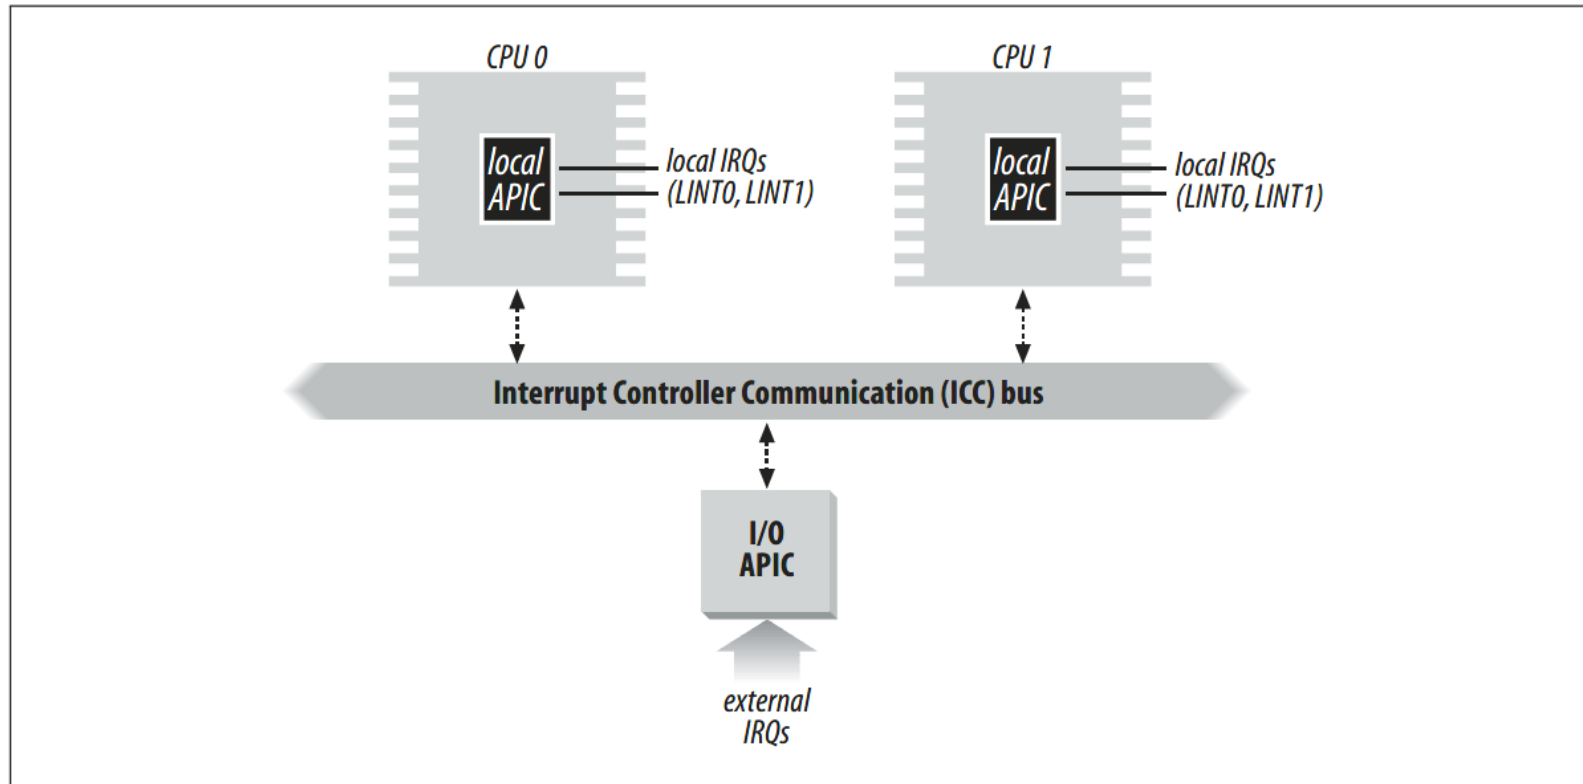
Interrupt Controller

Intel APIC(Advanced Programmable Interrupt Controller)

- 面向多核，每个CPU都有自己的local APIC
- 有一个I/O APIC负责接收所有设备信号，通过总线连接到所有核心的local APIC
 - 包含24个IRQ pin，以及一个24项的Interrupt Redirection Table
 - 优先级
 - 目标CPU
 - hardware IRQ到software IRQ的映射
 - 选择目标CPU的策略（动态/静态）
- 转发策略：
 - 静态：根据Redirection Table，可以一个/一些/广播
 - 动态：每个Local APIC有一个硬件task priority register，希望os每次reschedule的时候更新，这样每个dynamic IRQ到来时分发到最低优先级的CPU，优先级相同时RR
- cli时仅屏蔽local APIC中断
- mask一条IRQ line时会使I/O APIC不再响应这个irq，从而对所有CPU屏蔽掉一条IRQ
- 发送IPI：CPU将中断向量与目标Local APIC的标识符一起写入到自己Local APIC的Interrupt Command Register，从而触发IPI

Interrupt Controller

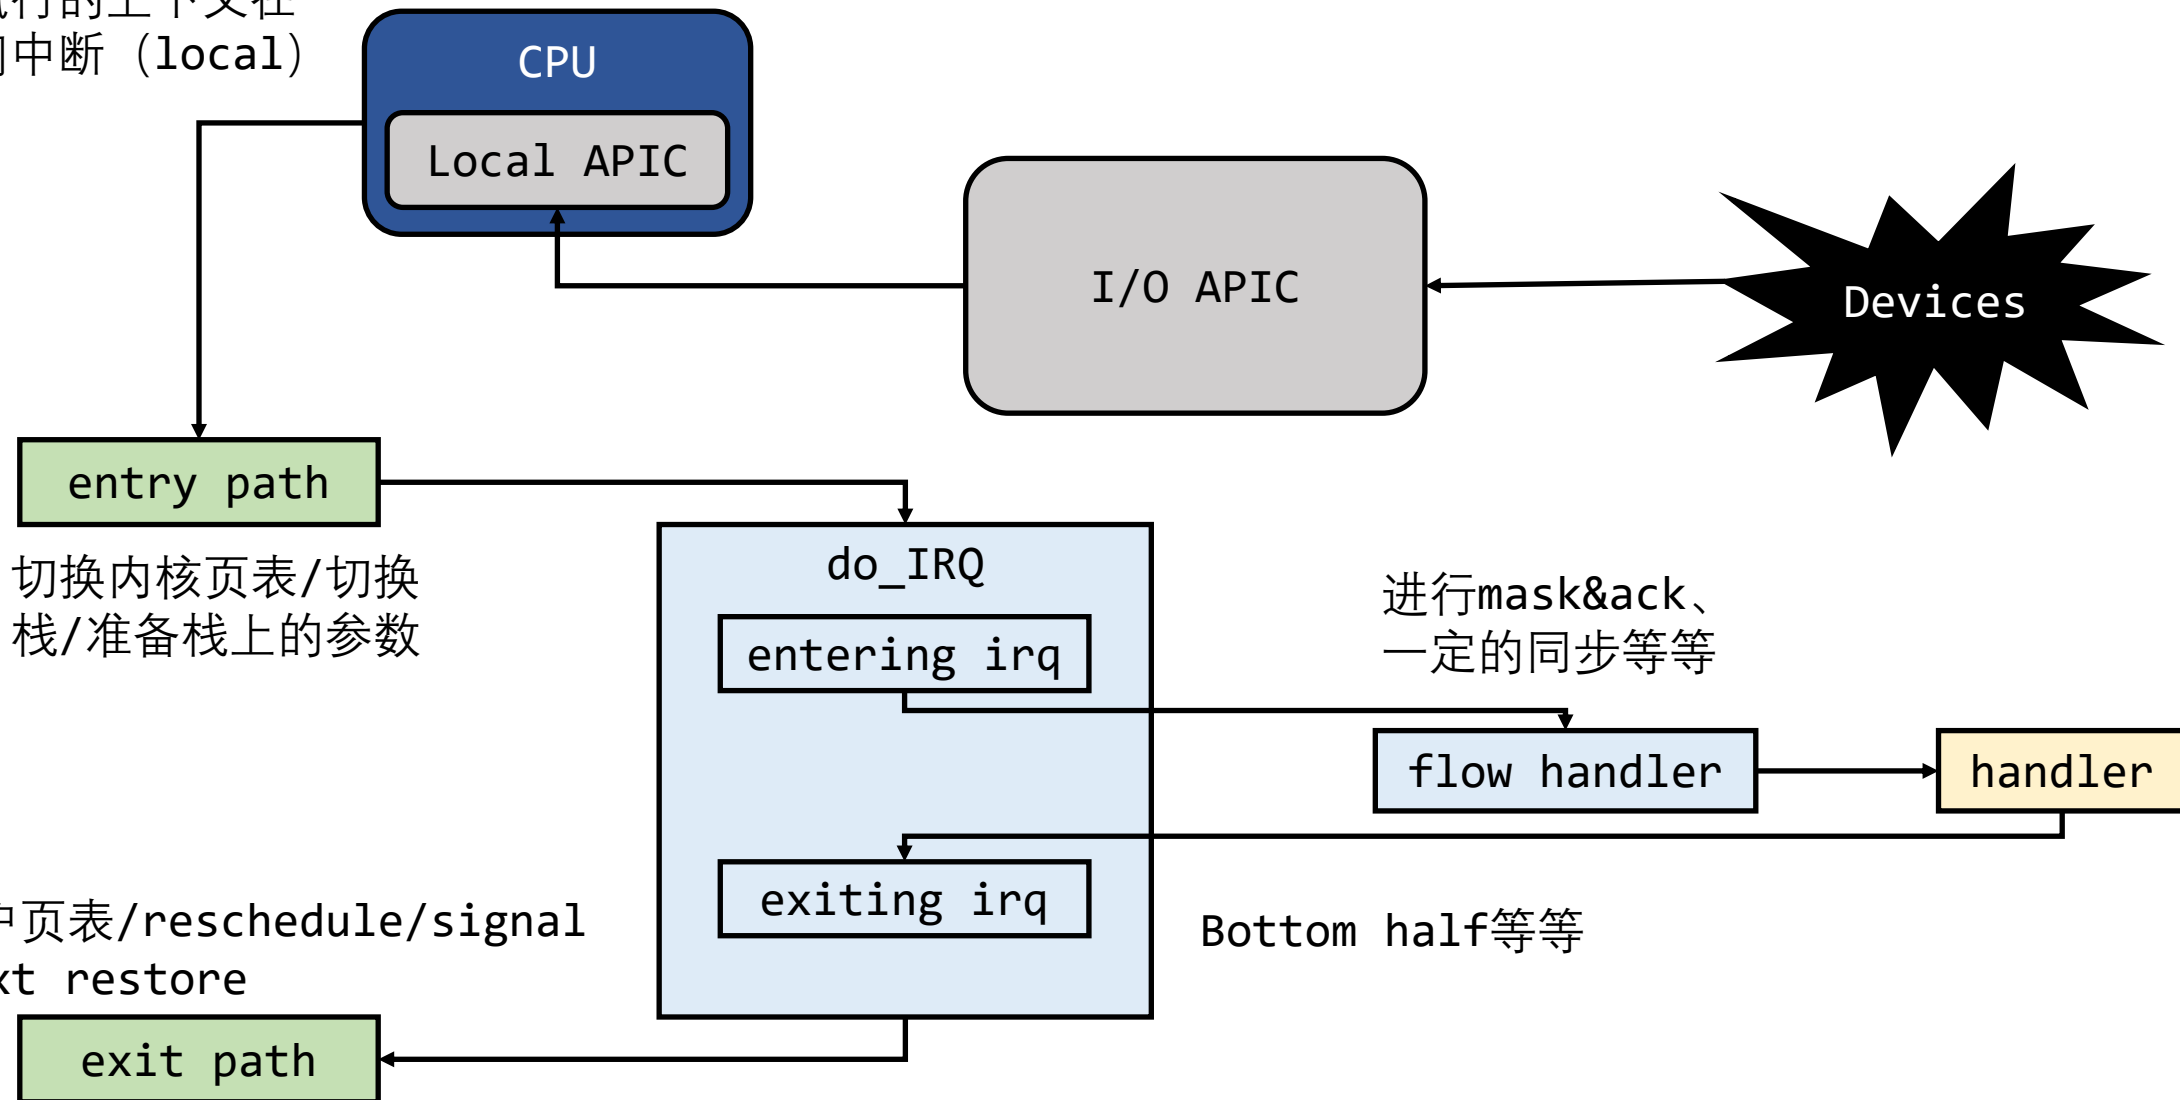
Intel APIC(Advanced Programmable Interrupt Controller)



Handling Overview

Handling Overview

保存当前执行的上下文在
栈上，关闭中断 (local)



Interrupt Context

Interrupt Context

- 发生中断时我们可能
 - 打断了其他用户程序
 - 打断了其他用户程序的内核path
 - 不管怎么样，都会有一个current进程，但这个current进程与我们无关
- 处于interrupt context时我们不能sleep或者阻塞，为什么？
 - interrupt disabled
 - interrupt enabled
 - 无辜的current
 - irq masked
- 这意味着
 - 我们甚至不能使用普通的内存分配（可能sleep）
 - 将不必要的内容推后到bottom half或者interrupt thread中

Interrupt Context

- handler要保证线程安全吗?
 - 如果handler只用于单一irq
 - 不需要, 因为执行handler时irq line会被mask, 其他核不会执行handler
 - 如果同一handler用于多个irq
 - 需要
- handler要保证可重入吗?
 - 如果handler只用于单一irq
 - 不需要, 因为执行handler时irq line会被mask, 当前核不会再重新进入handler
 - 如果同一handler用于多个irq
 - 执行handler时disable local interrupt
 - 不需要, 因为不会被打断, 没有机会执行重新进入irq handler
 - 执行handler时enable local interrupt
 - 需要

Handler Management

Handler Management

如何在某个irq上添加/移除我们的自定义handler呢？

Kernel API:

- request_irq/free_irq

```
int request_irq(  
    unsigned int irq,  
    irq_handler_t handler, // typedef int (*irq_handler_t)(irq, dev_id)  
    unsigned long irqflags, // 是否容忍共享irq line等等  
    const char *devname, // 会显示在/proc/interrupts下  
    void *dev_id); // 唯一标识符，一般传dev结构体指针const void *free_irq(  
    unsigned int irq,  
    void *dev_id);
```

Handler Management

```
struct irq_desc {  
    ...  
    struct irq_common_data {  
        ...  
        void *handler_data;  
        ...  
    } irq_common_data;  
    struct irq_data {  
        ...  
        struct irq_chip *chip;  
        void *chip_data;  
        ...  
    } irq_data;  
    irq_flow_handler_t handle_irq;  
    struct irq_action *action;  
    unsigned int istate;  
    struct depth;  
    const char *name;  
    ...  
}  
  
struct irq_chip {  
    ...  
    void (*irq_ack)(struct irq_data *data);  
    void (*irq_mask)(struct irq_data *data);  
    ...  
}  
  
struct irqaction {  
    irq_handler_t handler;  
    void *dev_id;  
    ...  
    struct irqaction *next;  
    ...  
    unsigned int irq;  
    unsigned int flags;  
    ...  
    const char *name;  
    ...  
}
```

The diagram illustrates the relationships between several kernel data structures in the context of interrupt handling:

- struct irq_desc**: Contains a **struct irq_common_data** (which points to **handler_data**), a **struct irq_data** (which points to **chip**), a **irq_flow_handler_t handle_irq**, a **struct irq_action *action** (which points to **irqaction**), **unsigned int istate**, **struct depth**, and **const char *name**.
- struct irq_chip**: Contains **void (*irq_ack)(struct irq_data *data)**, **void (*irq_mask)(struct irq_data *data)**, and other chip-specific functions.
- struct irqaction**: Contains **irq_handler_t handler**, **void *dev_id**, **struct irqaction *next**, **unsigned int irq**, **unsigned int flags**, **const char *name**, and other action-specific data.

Handler Management

```
struct irq_desc irq_desc[NR_IRQS] __cacheline_aligned_in_smp = {
    [0 ... NR_IRQS-1] = {
        .handle_irq = handle_bad_irq,
        .depth      = 1,
        .lock        = __RAW_SPIN_LOCK_UNLOCKED(irq_desc->lock),
    }
};
```


```
typedef struct irq_desc* vector_irq_t[NR_VECTORS];
DECLARE_PER_CPU(vector_irq_t, vector_irq);
```

```
DEFINE_PER_CPU(vector_irq_t, vector_irq) = {
    [0 ... NR_VECTORS - 1] = VECTOR_UNUSED,
};
```

Handling Overview

do_IRQ

保存当前执行的上下文在
栈上，关闭中断 (local)



do_IRQ

```
__visible unsigned int __irq_entry do_IRQ(struct pt_regs *regs)
{
    struct pt_regs *old_regs = set_irq_regs(regs);
    struct irq_desc *desc;
    /* high bit used in ret_from_code */
    unsigned vector = ~regs->orig_ax;

    entering_irq();

    /* entering_irq() tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    desc = __this_cpu_read(vector_irq[vector]);
    if (likely(!IS_ERR_OR_NULL(desc))) {
        if (IS_ENABLED(CONFIG_X86_32))
            handle_irq(desc, regs);
        else
            generic_handle_irq_desc(desc);
    } else {
        ack_APIC_irq();

        if (desc == VECTOR_UNUSED) {
            pr_emerg_ratelimited("%s: %d.%d No irq handler for vector\n",
                                __func__, smp_processor_id(),
                                vector);
        } else {
            __this_cpu_write(vector_irq[vector], VECTOR_UNUSED);
        }
    }

    exiting_irq();

    set_irq_regs(old_regs);
    return 1;
}
```

```
/*
 * Interrupt entry helper function.
 *
 * Entry runs with interrupts off. Stack layout at entry:
 * +-----+
 * | regs->ss
 * | regs->rsp
 * | regs->eflags
 * | regs->cs
 * | regs->ip
 * +-----+
 * | regs->orig_ax = ~(interrupt number)
 * +-----+
 * | return address
 * +-----+
 */
```

do_IRQ

```
__visible unsigned int __irq_entry do_IRQ(struct pt_regs *regs)
{
    struct pt_regs *old_regs = set_irq_regs(regs);
    struct irq_desc *desc;
    /* high bit used in ret_from_code */
    unsigned vector = ~regs->orig_ax;

    entering_irq();

    /* entering_irq() tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    desc = __this_cpu_read(vector_irq[vector]);
    if (likely(!IS_ERR_OR_NULL(desc))) {
        if (IS_ENABLED(CONFIG_X86_32))
            handle_irq(desc, regs);
        else
            generic_handle_irq_desc(desc);
    } else {
        ack_APIC_irq();

        if (desc == VECTOR_UNUSED) {
            pr_emerg_ratelimited("%s: %d.%d No irq handler for vector\n",
                                __func__, smp_processor_id(),
                                vector);
        } else {
            __this_cpu_write(vector_irq[vector], VECTOR_UNUSED);
        }
    }

    exiting_irq();

    set_irq_regs(old_regs);
    return 1;
}
```

1. save/restore irq context

do_IRQ

```
__visible unsigned int __irq_entry do_IRQ(struct pt_regs *regs)
{
    struct pt_regs *old_regs = set_irq_regs(regs);
    struct irq_desc *desc;
    /* high bit used in ret_from_code */
    unsigned vector = ~regs->orig_ax;

    entering_irq();

    /* entering_irq() tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    desc = __this_cpu_read(vector_irq[vector]);
    if (unlikely(!IS_ERR_OR_NULL(desc))) {
        if (IS_ENABLED(CONFIG_X86_32))
            handle_irq(desc, regs);
        else
            generic_handle_irq_desc(desc);
    } else {
        ack_APIC_irq();

        if (desc == VECTOR_UNUSED) {
            pr_emerg_ratelimited("%s: %d.%d No irq handler for vector\n",
                                __func__, smp_processor_id(),
                                vector);
        } else {
            __this_cpu_write(vector_irq[vector], VECTOR_UNUSED);
        }
    }

    exiting_irq();

    set_irq_regs(old_regs);
    return 1;
}
```

2. 获取中断描述符

do_IRQ

```
__visible unsigned int __irq_entry do_IRQ(struct pt_regs *regs)
{
    struct pt_regs *old_regs = set_irq_regs(regs);
    struct irq_desc *desc;
    /* high bit used in ret_from_code */
    unsigned vector = ~regs->orig_ax;

    entering_irq();

    /* entering_irq() tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    desc = __this_cpu_read(vector_irq[vector]);
    if (likely(!IS_ERR_OR_NULL(desc))) {
        if (IS_ENABLED(CONFIG_X86_32))
            handle_irq(desc, regs);
        else
            generic_handle_irq_desc(desc);
    } else {
        ack_APIC_irq();

        if (desc == VECTOR_UNUSED) {
            pr_emerg_ratelimited("%s: %d.%d No irq handler for vector\n",
                                __func__, smp_processor_id(),
                                vector);
        } else {
            __this_cpu_write(vector_irq[vector], VECTOR_UNUSED);
        }
    }

    exiting_irq();

    set_irq_regs(old_regs);
    return 1;
}
```

3. 32/64-bit handling

```
/*
 * Architectures call this to let the generic IRQ layer
 * handle an interrupt.
 */
static inline void generic_handle_irq_desc(struct irq_desc *desc)
{
    desc->handle_irq(desc);
}
```

irq_flow_handler_t

```
struct irq_desc {  
    ...  
    struct irq_common_data {  
        ...  
        void *handler_data;  
        ...  
    } irq_common_data;  
    struct irq_data {  
        ...  
        struct irq_chip *chip;  
        void *chip_data;  
        ...  
    } irq_data;  
    irq_flow_handler_t handle_irq;  
    struct irq_action *action;  
    unsigned int istate;  
    struct depth;  
    const char *name;  
    ...  
}  
  
struct irq_chip {  
    ...  
    void (*irq_ack)(struct irq_data *data);  
    void (*irq_mask)(struct irq_data *data);  
    ...  
}  
  
struct irqaction {  
    irq_handler_t handler;  
    void *dev_id;  
    ...  
    struct irqaction *next;  
    ...  
    unsigned int irq;  
    unsigned int flags;  
    ...  
    const char *name;  
    ...  
}
```

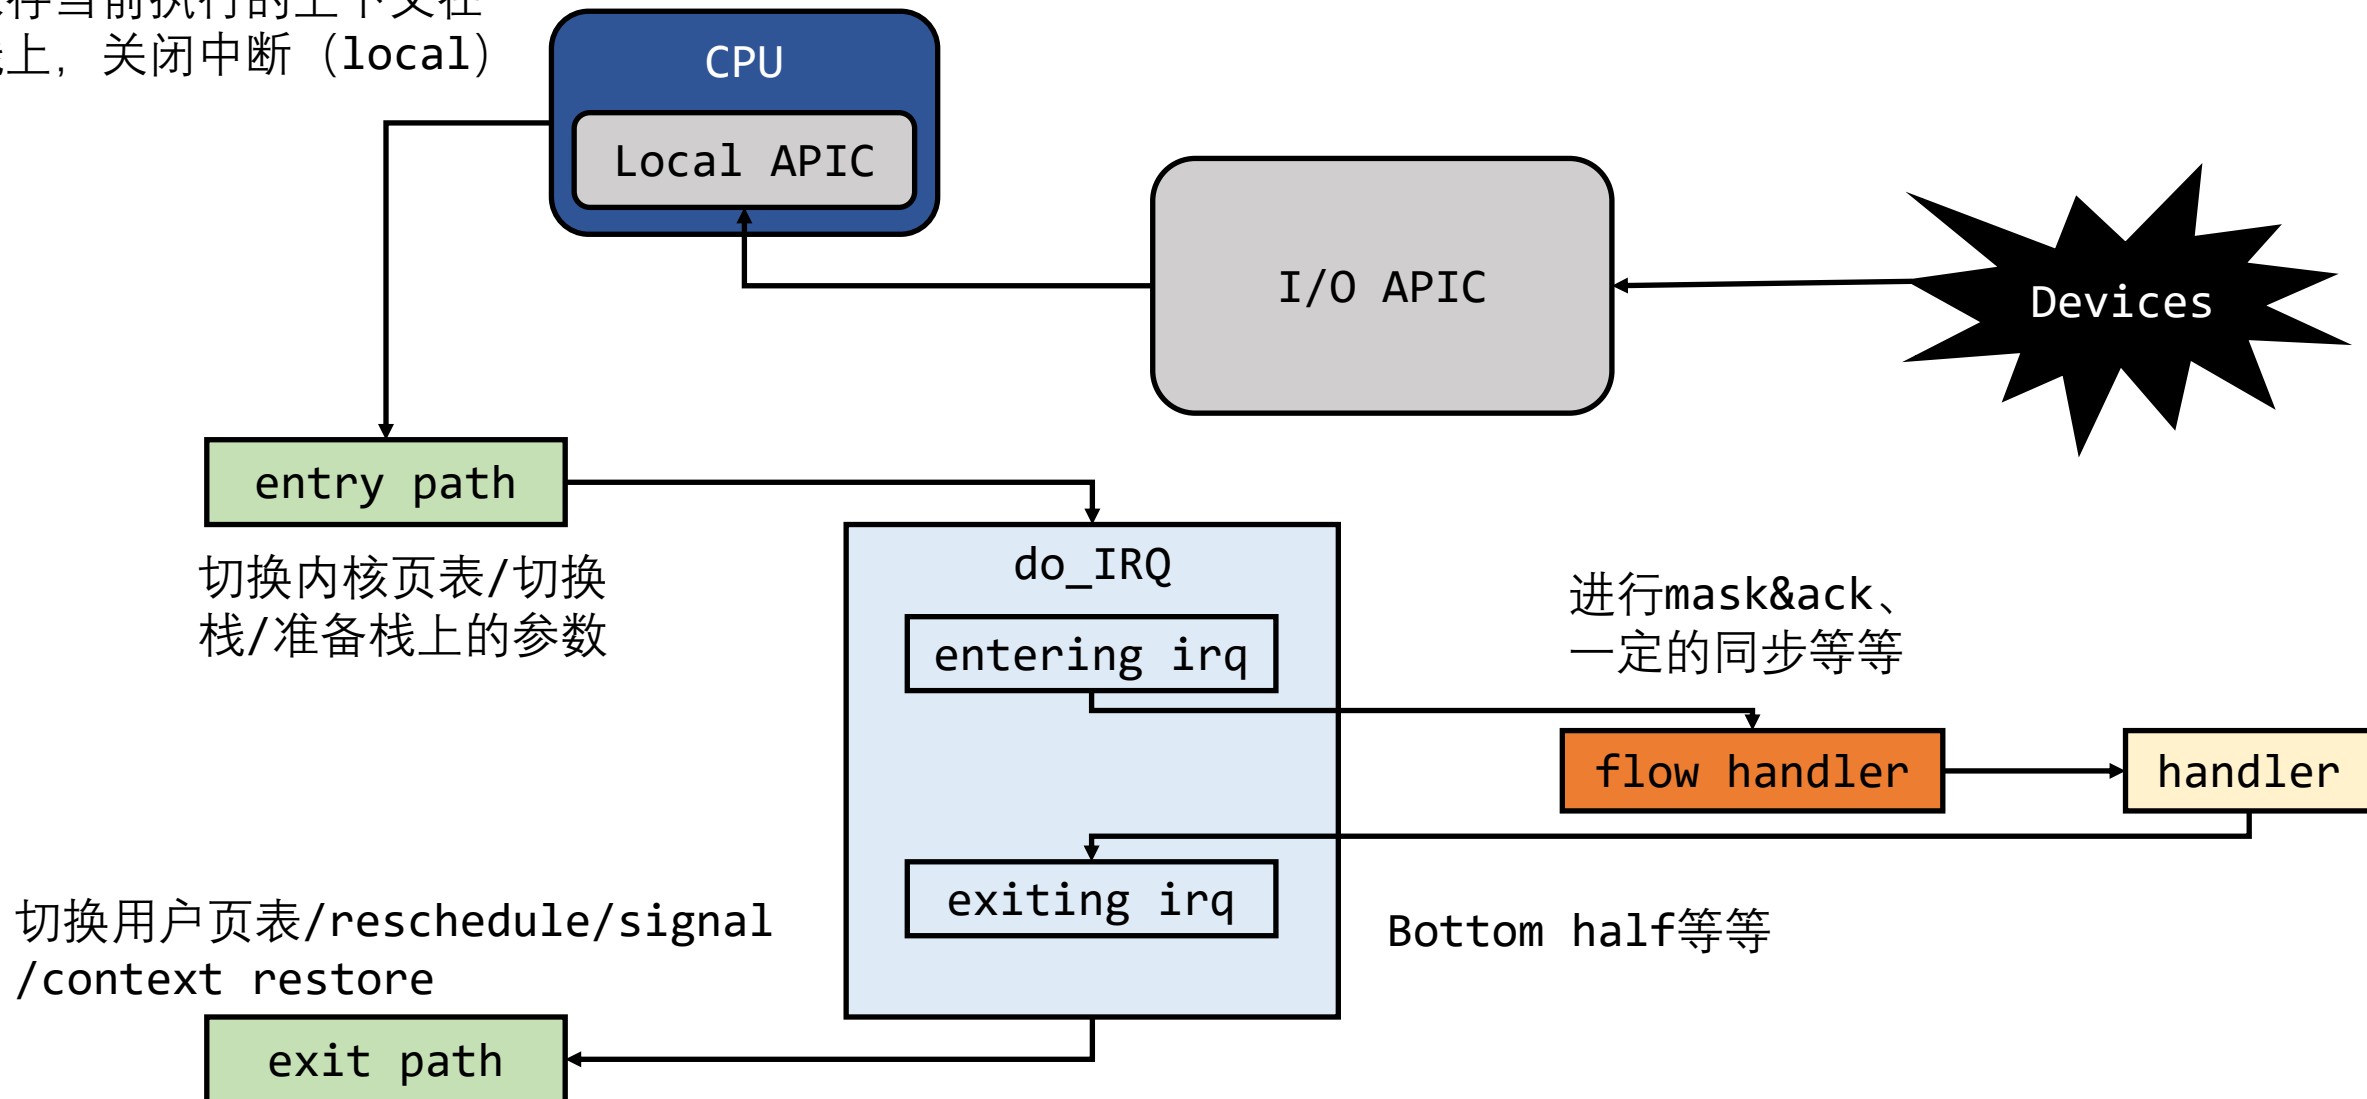
The diagram illustrates the relationships between kernel data structures for interrupt handling. Blue arrows indicate the following connections:

- From `handler_data` in `struct irq_common_data` to `irq_chip`.
- From `chip` in `struct irq_data` to `irq_chip`.
- From `action` in `struct irq_desc` to `irqaction`.

A red box highlights the `irq_flow_handler_t handle_irq;` member in `struct irq_desc`.

irq_flow_handler_t

保存当前执行的上下文在
栈上，关闭中断 (local)



irq_flow_handler_t

interrupt handler types:

- Level type
- Edge type
- Simple type
- Fast EOI type
- Per CPU type

handle_level_irq

```
void handle_level_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);
    mask_ack_irq(desc);

    if (!irq_may_run(desc))
        goto out_unlock;

    ...

    handle_irq_event(desc);

    ...

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```


handle_level_irq

```
irqreturn_t handle_irq_event(struct irq_desc *desc)
{
    irqreturn_t ret;

    desc->istate &= ~IRQS_PENDING;
    irqd_set(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    raw_spin_unlock(&desc->lock);

    ret = handle_irq_event_percpu(desc);

    raw_spin_lock(&desc->lock);
    irqd_clear(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    return ret;
}
```

- 标志自己正在处理，release锁
- handle irq
- 重新hold锁，抹去自己正在处理的标志

handle_level_irq

```
irqreturn_t __handle_irq_event_percpu(struct irq_desc *desc, unsigned int *flags)
{
    irqreturn_t retval = IRQ_NONE;
    unsigned int irq = desc->irq_data.irq;
    struct irqaction *action;

    record_irq_time(desc);

    for_each_action_of_desc(desc, action) {
        irqreturn_t res;

        trace_irq_handler_entry(irq, action);
        res = action->handler(irq, action->dev_id);
        trace_irq_handler_exit(irq, action, res);

        if (WARN_ONCE(!irqs_disabled(), "irq %u handler %pS enabled interrupts\n",
                    irq, action->handler))
            local_irq_disable();

        switch (res) {
        case IRQ_WAKE_THREAD:
            /*
             * Catch drivers which return WAKE_THREAD but
             * did not set up a thread function
             */
            if (unlikely(!action->thread_fn)) {
                warn_no_thread(irq, action);
                break;
            }

            __irq_wake_thread(desc, action);

            /* Fall through - to add to randomness */
        case IRQ_HANDLED:
            *flags |= action->flags;
            break;

        default:
            break;
        }

        retval |= res;
    }
}
```

遍历所有action（也即所有的shared handler）
尝试响应irq

通过检查返回值可知中断是否有handler响应

handle_level_irq

```
void handle_level_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);
    mask_ack_irq(desc);

    if (!irq_may_run(desc))
        goto out_unlock;

    ...

    handle_irq_event(desc);

    ...

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

```
static bool irq_may_run(struct irq_desc *desc)
{
    unsigned int mask = IRQD_IRQ_INPROGRESS | IRQD_WAKEUP_ARMED;

    if (!irqd_has_set(&desc->irq_data, mask))
        return true;

    ...

}
```

```
irqreturn_t handle_irq_event(struct irq_desc *desc)
{
    irqreturn_t ret;

    desc->istate &= ~IRQS_PENDING;
    irqd_set(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    raw_spin_unlock(&desc->lock);

    ret = handle_irq_event_percpu(desc);

    raw_spin_lock(&desc->lock);
    irqd_clear(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    return ret;
}
```

handle_edge_irq

```
void handle_edge_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);

    ...

    if (!irq_may_run(desc)) {
        desc->istate |= IRQS_PENDING;
        mask_ack_irq(desc);
        goto out_unlock;
    }

    ...

    /* Start handling the irq */
    desc->irq_data.chip->irq_ack(&desc->irq_data);

    do {

        ...

        if (unlikely(desc->istate & IRQS_PENDING)) {
            if (!irqd_irq_disabled(&desc->irq_data) &&
                irqd_irq_masked(&desc->irq_data))
                unmask_irq(desc);
        }

        handle_irq_event(desc);

    } while ((desc->istate & IRQS_PENDING) &&
            !irqd_irq_disabled(&desc->irq_data));

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

handle_edge_irq

core1

```
void handle_edge_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);

    ...

    /* Start handling the irq */
    desc->irq_data.chip->irq_ack(&desc->irq_data);

    do {

        ...

        if (unlikely(desc->istate & IRQS_PENDING)) {
            if (!irqd_irq_disabled(&desc->irq_data) &&
                irqd_irq_masked(&desc->irq_data))
                unmask_irq(desc);
        }

        handle_irq_event(desc);

    } while ((desc->istate & IRQS_PENDING) &&
            !irqd_irq_disabled(&desc->irq_data));

    out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

1. ack(而不mask)

2. handle irq

3. 检查pending, 处理新来的edge中断 (因为中间release了锁, 所以可能其他核会重入这部分代码设置pending)

```
irqreturn_t handle_irq_event(struct irq_desc *desc)
{
    irqreturn_t ret;

    desc->istate &= ~IRQS_PENDING;
    irqd_set(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    raw_spin_unlock(&desc->lock);

    ret = handle_irq_event_percpu(desc);

    raw_spin_lock(&desc->lock);
    irqd_clear(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    return ret;
}
```

handle_edge_irq

core2

```
void handle_edge_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);

    ...

    if (!irq_may_run(desc)) {
        desc->istate |= IRQS_PENDING;
        mask_ack_irq(desc);
        goto out_unlock;
    }

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

```
static bool irq_may_run(struct irq_desc *desc)
{
    unsigned int mask = IRQD_IRQ_INPROGRESS | IRQD_WAKEUP_ARMED;

    if (!irqd_has_set(&desc->irq_data, mask))
        return true;

    ...

}
```

标记pending（用于core 1的loop检测，告诉他处理新来的edge中断）

mask & ack：所有核不再接收这条irq上的中断信号，最多只能pending一个edge中断

handle_edge_irq

core1

```
void handle_edge_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);

    ...

    /* Start handling the irq */
    desc->irq_data.chip->irq_ack(&desc->irq_data);

    do {

        ...

        if (unlikely(desc->istate & IRQS_PENDING)) {
            if (!irqd_irq_disabled(&desc->irq_data) &&
                irqd_irq_masked(&desc->irq_data))
                unmask_irq(desc);
        }

        handle_irq_event(desc);

    } while ((desc->istate & IRQS_PENDING) &&
            !irqd_irq_disabled(&desc->irq_data));

    out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

1. ack(而不mask)

4. 取出pending,
重新unmask

2. handle irq

3. 检查pending, 处理新来的edge中断 (因为中间release了锁, 所以可能其他核会重入这部分代码设置pending)

```
irqreturn_t handle_irq_event(struct irq_desc *desc)
{
    irqreturn_t ret;

    desc->istate &= ~IRQS_PENDING;
    irqd_set(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    raw_spin_unlock(&desc->lock);

    ret = handle_irq_event_percpu(desc);

    raw_spin_lock(&desc->lock);
    irqd_clear(&desc->irq_data, IRQD_IRQ_INPROGRESS);
    return ret;
}
```

handle_edge_irq

```
void handle_edge_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);

    ...

    if (!irq_may_run(desc)) {
        desc->istate |= IRQS_PENDING;
        mask_ack_irq(desc);
        goto out_unlock;
    }

    ...

    /* Start handling the irq */
    desc->irq_data.chip->irq_ack(&desc->irq_data);

    do {

        ...

        if (unlikely(desc->istate & IRQS_PENDING)) {
            if (!irqd_irq_disabled(&desc->irq_data) &&
                irqd_irq_masked(&desc->irq_data))
                unmask_irq(desc);
        }

        handle_irq_event(desc);

    } while ((desc->istate & IRQS_PENDING) &&
        !irqd_irq_disabled(&desc->irq_data));

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```

```
void handle_level_irq(struct irq_desc *desc)
{
    raw_spin_lock(&desc->lock);
    mask_ack_irq(desc);

    if (!irq_may_run(desc))
        goto out_unlock;

    ...

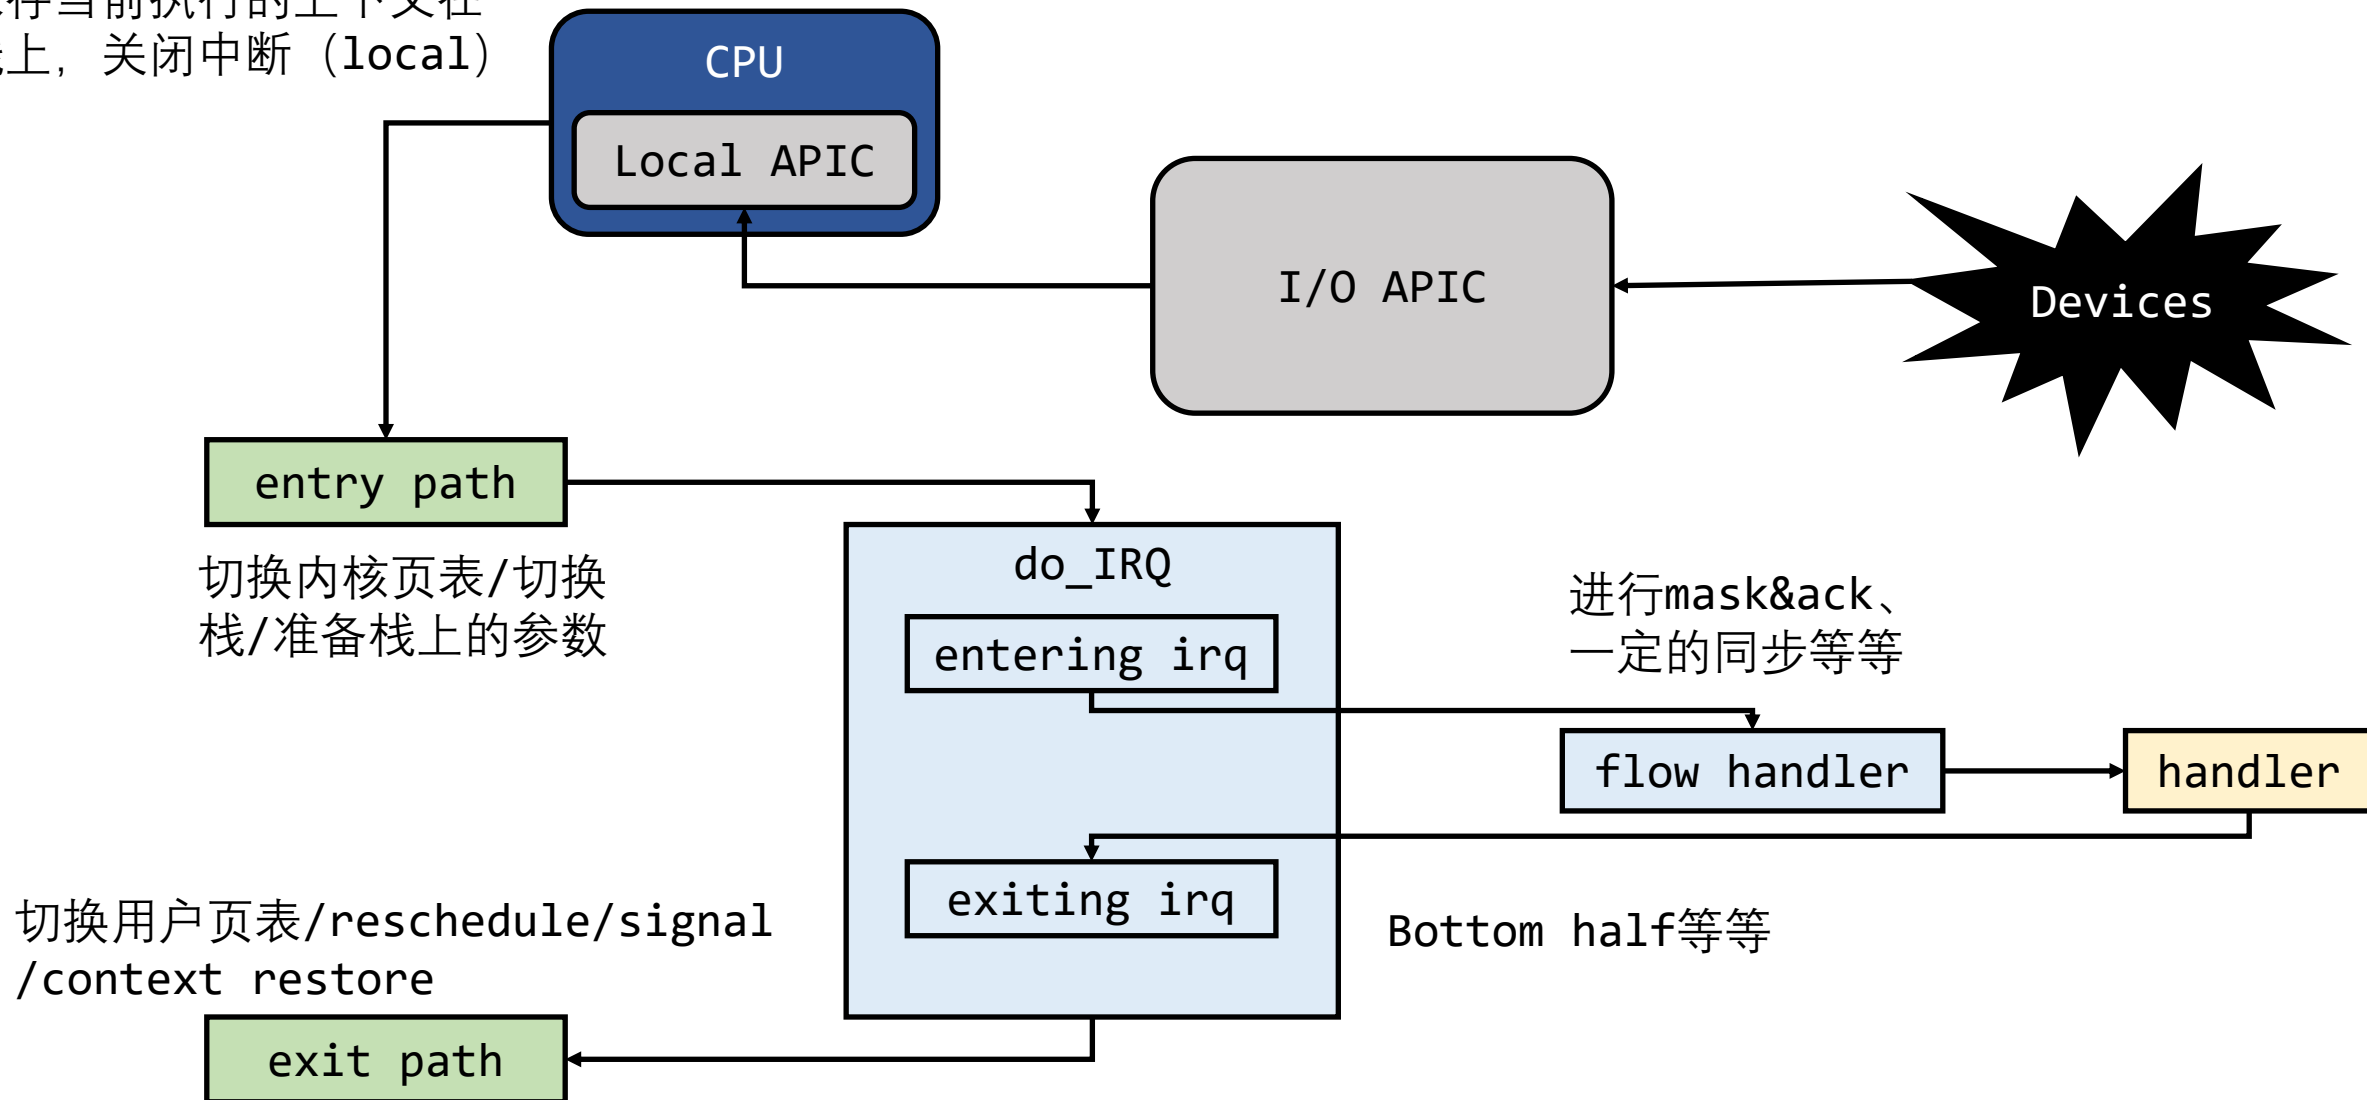
    handle_irq_event(desc);

    ...

out_unlock:
    raw_spin_unlock(&desc->lock);
}
```


irq_flow_handler_t

保存当前执行的上下文在
栈上，关闭中断 (local)



do_IRQ

```
__visible unsigned int __irq_entry do_IRQ(struct pt_regs *regs)
{
    struct pt_regs *old_regs = set_irq_regs(regs);
    struct irq_desc *desc;
    /* high bit used in ret_from_code */
    unsigned vector = ~regs->orig_ax;

    entering_irq();

    /* entering_irq() tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    desc = __this_cpu_read(vector_irq[vector]);
    if (likely(!IS_ERR_OR_NULL(desc))) {
        if (IS_ENABLED(CONFIG_X86_32))
            handle_irq(desc, regs);
        else
            generic_handle_irq_desc(desc);
    } else {
        ack_APIC_irq();

        if (desc == VECTOR_UNUSED) {
            pr_emerg_ratelimited("%s: %d.%d No irq handler for vector\n",
                                __func__, smp_processor_id(),
                                vector);
        } else {
            __this_cpu_write(vector_irq[vector], VECTOR_UNUSED);
        }
    }

    exiting_irq();

    set_irq_regs(old_regs);
    return 1;
}
```

3. 32/64-bit handling

```
/*
 * Architectures call this to let the generic IRQ layer
 * handle an interrupt.
 */
static inline void generic_handle_irq_desc(struct irq_desc *desc)
{
    desc->handle_irq(desc);
}
```

handle_irq

```
void handle_irq(struct irq_desc *desc, struct pt_regs *regs)
{
    int overflow = check_stack_overflow();

    if (user_mode(regs) || !execute_on_irq_stack(overflow, desc)) {
        if (unlikely(overflow))
            print_stack_overflow();
        generic_handle_irq_desc(desc);
    }
}
```

```
static inline int execute_on_irq_stack(int overflow, struct irq_desc *desc)
{
    struct irq_stack *curstk, *irqstk;
    u32 *isp, *prev_esp, arg1;

    curstk = (struct irq_stack *) current_stack();
    irqstk = __this_cpu_read(hardirq_stack_ptr);

    /*
     * this is where we switch to the IRQ stack. However, if we are
     * already using the IRQ stack (because we interrupted a hardirq
     * handler) we can't do that and just have to keep using the
     * current stack (which is the irq stack already after all)
     */
    if (unlikely(curstk == irqstk))
        return 0;

    isp = (u32 *) ((char *)irqstk + sizeof(*irqstk));

    /* Save the next esp at the bottom of the stack */
    prev_esp = (u32 *)irqstk;
    *prev_esp = current_stack_pointer;

    if (unlikely(overflow))
        call_on_stack(print_stack_overflow, isp);

    asm volatile("xchgl %%ebx,%%esp \n"
                 CALL_NOSPEC
                 "movl %%ebx,%%esp \n"
                 : "=a" (arg1), "=b" (isp)
                 : "0" (desc), "1" (isp),
                 [thunk_target] "D" (desc->handle_irq)
                 : "memory", "cc", "ecx");
    return 1;
}
```

Nested Interrupt

IRQF_DISABLED

```
irqreturn_t __handle_irq_event_percpu(struct irq_desc *desc, unsigned int *flags)
{
    irqreturn_t retval = IRQ_NONE;
    unsigned int irq = desc->irq_data.irq;
    struct irqaction *action;

    record_irq_time(desc);

    for_each_action_of_desc(desc, action) {
        irqreturn_t res;

        trace_irq_handler_entry(irq, action);
        res = action->handler(irq, action->dev_id);
        trace_irq_handler_exit(irq, action, res);

        if (WARN_ONCE(!irqs_disabled(), "irq %u handler %pS enabled interrupts\n",
                    irq, action->handler))
            local_irq_disable();

        switch (res) {
        case IRQ_WAKE_THREAD:
            /*
             * Catch drivers which return WAKE_THREAD but
             * did not set up a thread function
             */
            if (unlikely(!action->thread_fn)) {
                warn_no_thread(irq, action);
                break;
            }

            __irq_wake_thread(desc, action);

            /* Fall through - to add to randomness */
        case IRQ_HANDLED:
            *flags |= action->flags;
            break;

        default:
            break;
        }

        retval |= res;
    }
}
```

遍历所有action（也即所有的shared handler）
尝试响应irq

```
if (!(action->flags & IRQF_DISABLED))
    local_irq_enable_in_hard_irq();
```

通过检查返回值可知中断是否有handler响应

Nested Interrupt

移除IRQF_DISABLED

- 世界不是完美的，我们当然希望所有的irq handler都能够快速执行critical任务with interrupt disabled，但限于有些老旧的硬件以及开发handler人员的水平，很多handler执行是比较慢的。
- 举个例子，IDE硬盘的数据传输相比串口的UART就是慢的，由于UART只能保存一个字符，我们应该尽可能地让UART的先执行，避免过多的字符丢失。从程序上体现就是嵌套中断。

Nested Interrupt

那么为什么又要重新提出（默认）关闭嵌套中断呢？

- 设备硬件、cpu的性能都比以前高了很多（所以所有handler的执行、数据传输都更快了）
- 内核的bottom half机制日趋完善、稳定，很多工作都可以推迟到bottom half中，慢中断（允许嵌套中断）与快中断之间的界限没那么明显了
- 而默认允许嵌套中断也带来了一系列问题
 - 发生嵌套时对cache不友好
 - handler的运行时间不稳定（被嵌套打断）
 - 中断栈的溢出（即使切换到中断栈上仍然可能会溢出）

Nested Interrupt

导火索：针对中断栈溢出的一个patch

- commit的人对内核做出修改，当执行嵌套中断的时候如果发现中断栈被用了超过一半，就不管是否允许中断打开而直接关着中断执行handler
- 但是这个人也表示，这其实治标不治本，核心问题还是默认的嵌套中断：

如果一个handler写得足够好，很快就能执行完，那么执行它的时候非要开着中断（别的中断会打断、嵌套）有啥意义呢？

Nested Interrupt

一开始Linux并不赞同取消默认的嵌套中断，一些例子：

- 一些handler注定要执行很多工作，而且这是很难改变的，比如很多挑剔的嵌入式硬件以及它们孱弱的处理器
- 有些handler不开着中断没法正常运行，比如有的handler可能要循环等待jiffies被增加多少（由timer interrupt更新），关掉中断就成死循环了
- 还有一些handler可能对延迟容忍度很低，必须立马得到执行（以嵌套的方式）而不是等着前面的handler执行完（重新打开中断）

Nested Interrupt

所以一个问题是，如果我们取消了默认的嵌套中断机制，系统还能正常运行吗？

- 一个很有趣的现象是，很多内核开着lockdep checker已经运行很多年了，而它会使得handler运行在interrupt disabled状态下
- 还有就是动态tick部分的代码会在系统空闲的时候直接disable时钟中断，这也导致jiffies无法得到更新，因此很多驱动针对jiffies不更新已经做出了改善
- 另外就是realtime tree的开发者花很多精力在中断处理的延迟问题上，超时是其中最不可接收的问题之一，所以它们的代码也会针对中断关闭的问题及时做出更新
- threaded interrupt handler模块的支持使得现有的驱动代码可以很好地迁移，直接移动到handler的thread中去即可（thread中中断是打开的）

Nested Interrupt

当然，仍然还有问题等待未来的解决

- 对嵌入系统来说，handler的thread带来的响应延迟是很难接受的
- 对于实在需要打开中断的，可以借助`local_irq_enable_in_hardirq`在handler内部打开中断

<https://lwn.net/Articles/380931/>

感谢观看

■ 彭天祥