

M³x: Autonomous Accelerators via Context-Enabled Fast-Path Communication

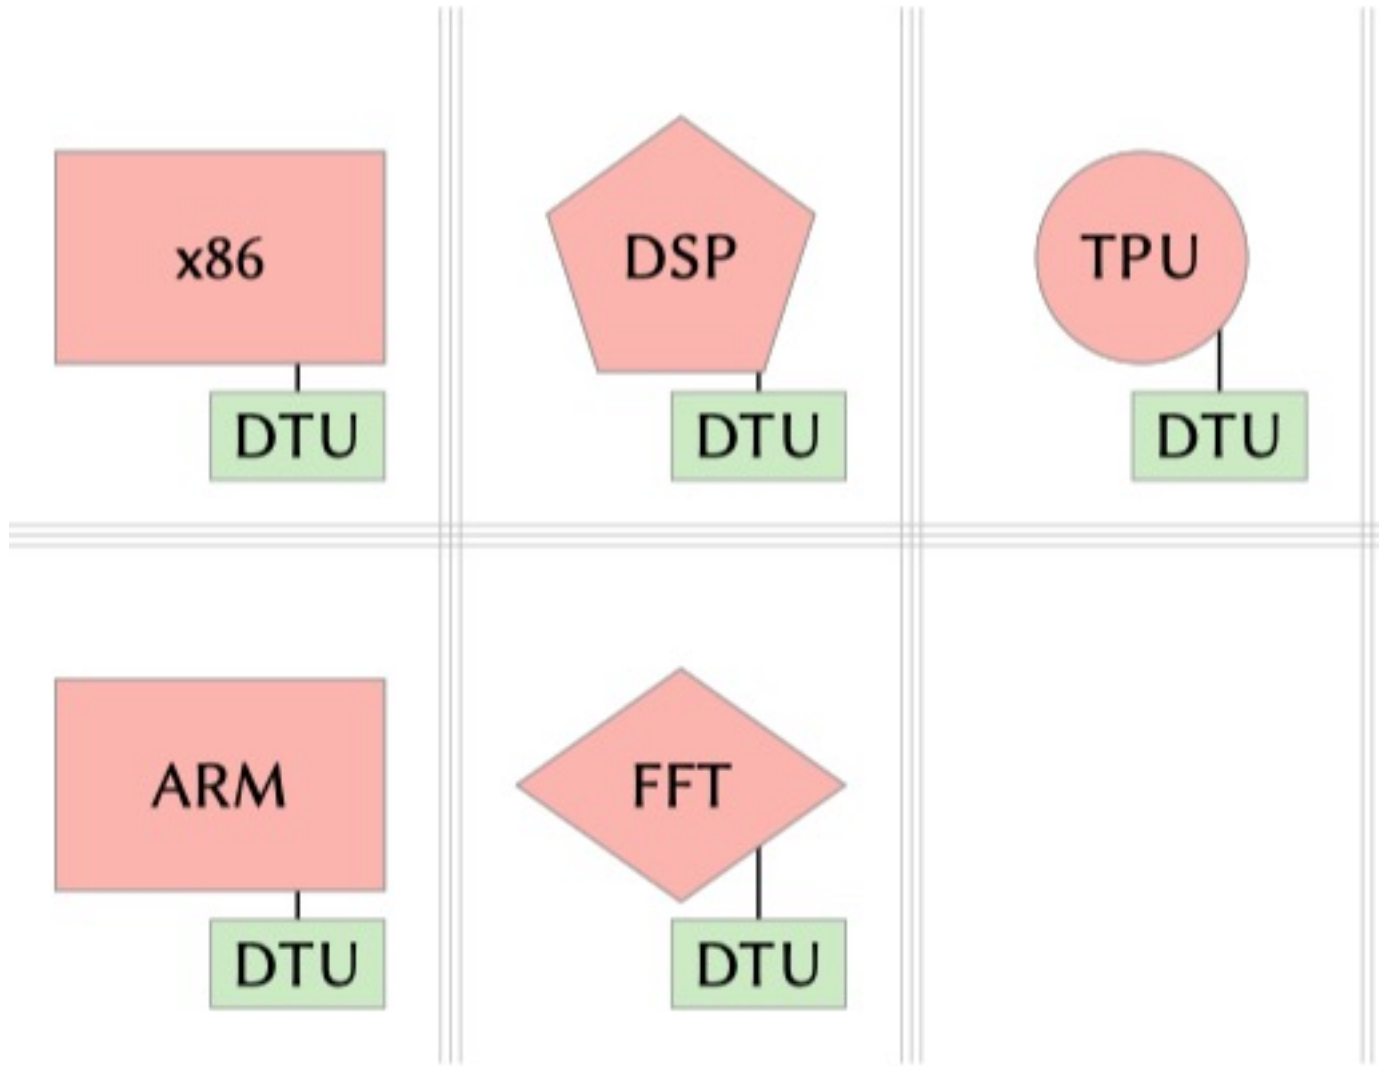
Contributions

- Rethink system architecture based on M^3
 - Hardware/operating-system co-design for heterogeneous systems
 - Simulation based on gem5
- Not built upon cache coherency
 - Costs (area, power, complexity, performance) increases with system size
 - More challenging for heterogeneous systems
 - Unclear whether future systems will be (globally) coherent
- Focus on fixed-function accelerators
 - Most difficult to support as “first-class citizens”
 - Provide none of the features OSes need

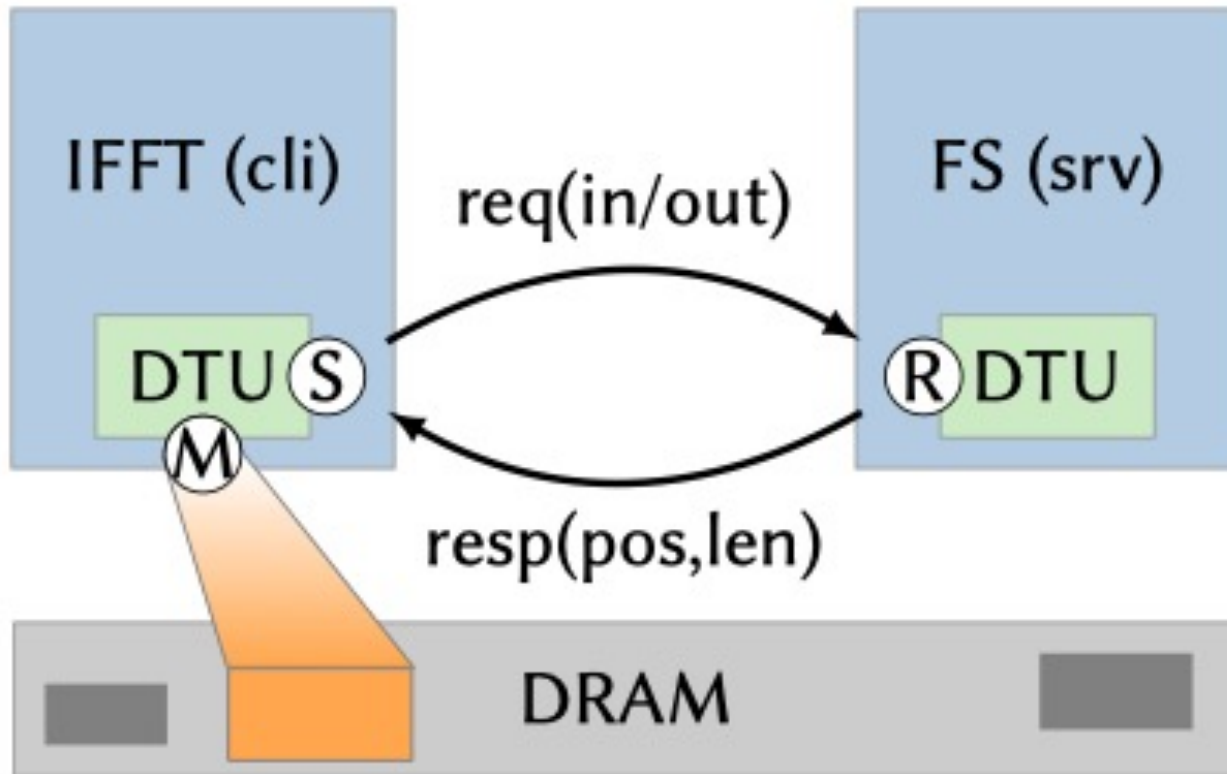
Outline

- System Architecture and Background
- Autonomous Stream-Processing Accelerators
- Fast-Path Communication vs. Context Switching

System Architecture and Background




- Key Ideas of M³[1]:
 - DTU as uniform interface
 - Kernel controls user tiles remotely



Autonomous Stream-Processing Accelerators

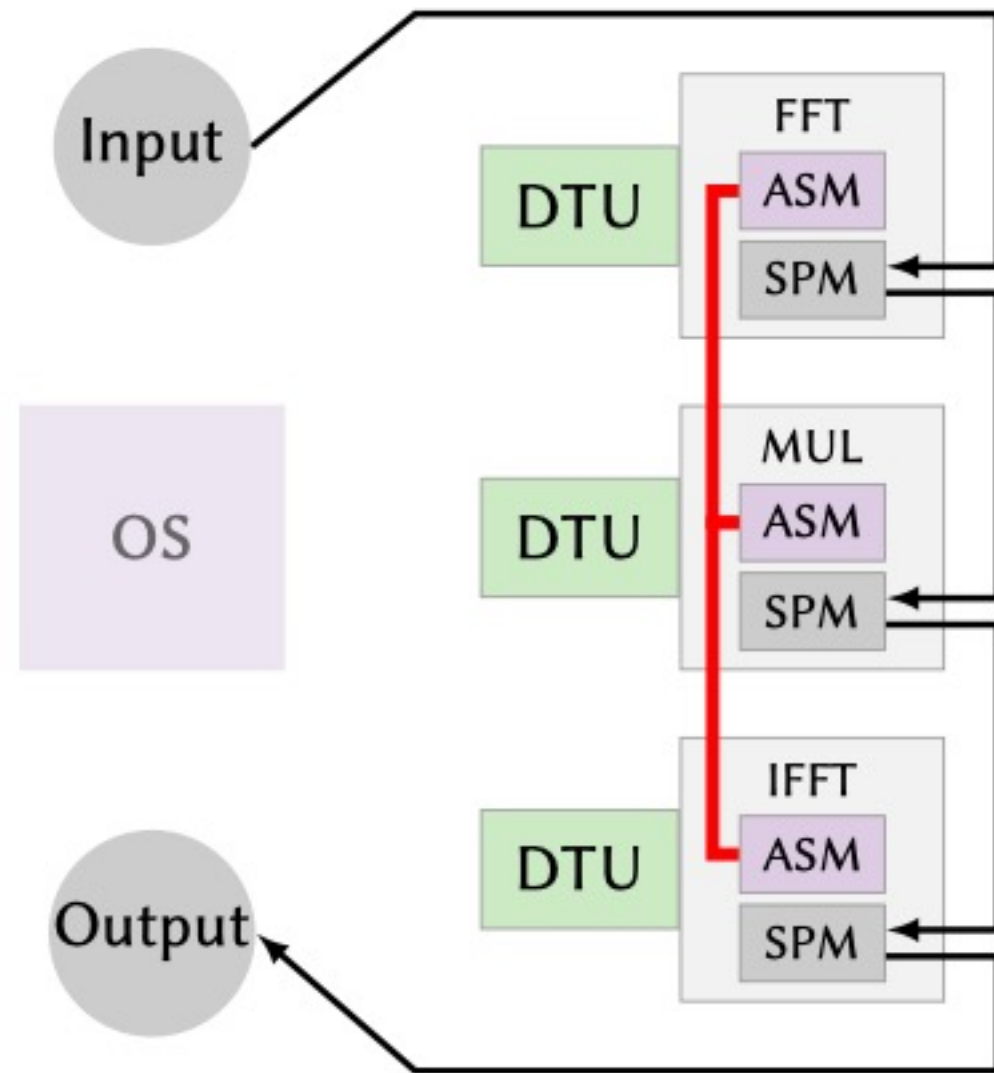
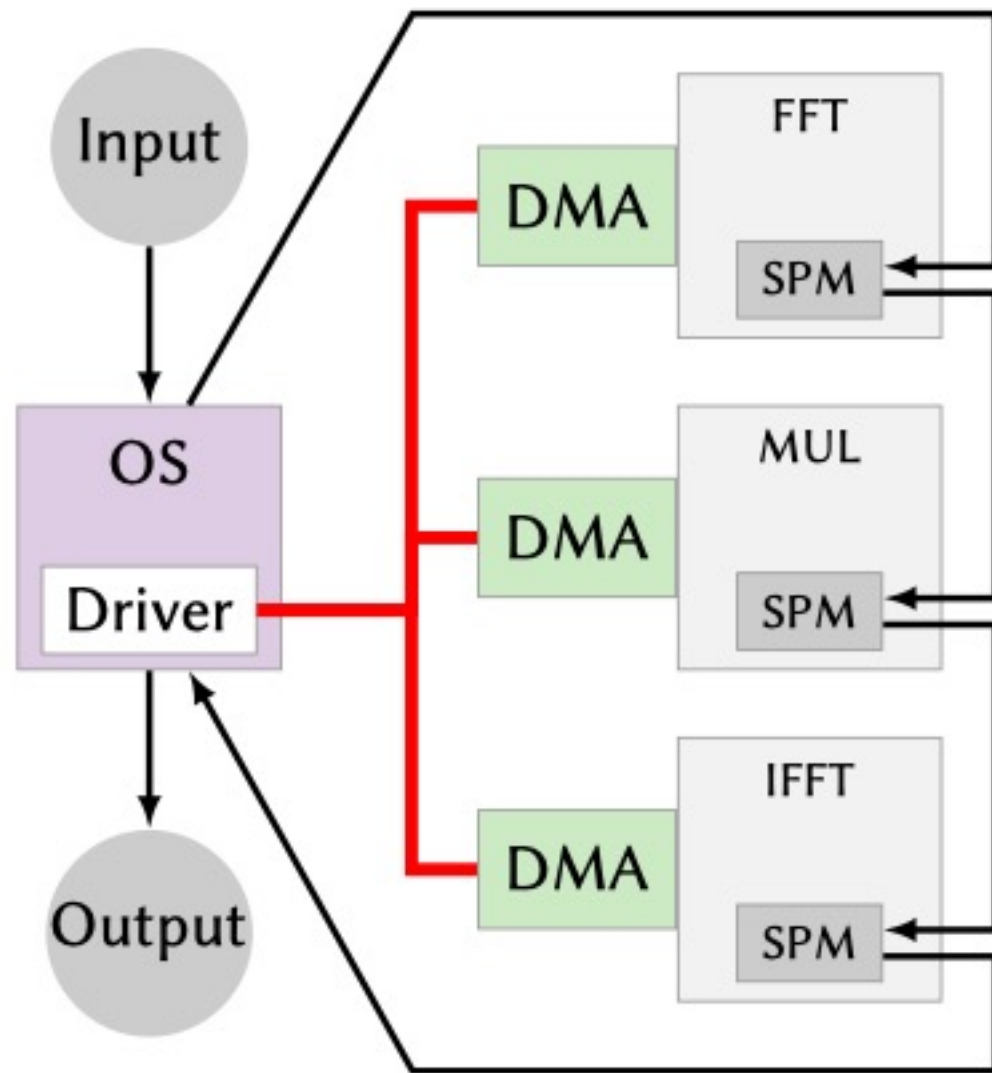
- File protocol
 - Data in memory
 - Msg channel between client and server
 - Req (in) for next input piece
 - Req (out) for next output piece
 - Server configures client's memory EP
 - Client accesses data via DTU
 - Used by all CUs




Autonomous Stream-Processing Accelerators

- Off-the-shelf accelerators
- Accelerator Support Module (ASM):
 - Interacts with DTU and accelerator
 - Implements file protocol for input and output channel


Assisted vs. Autonomous



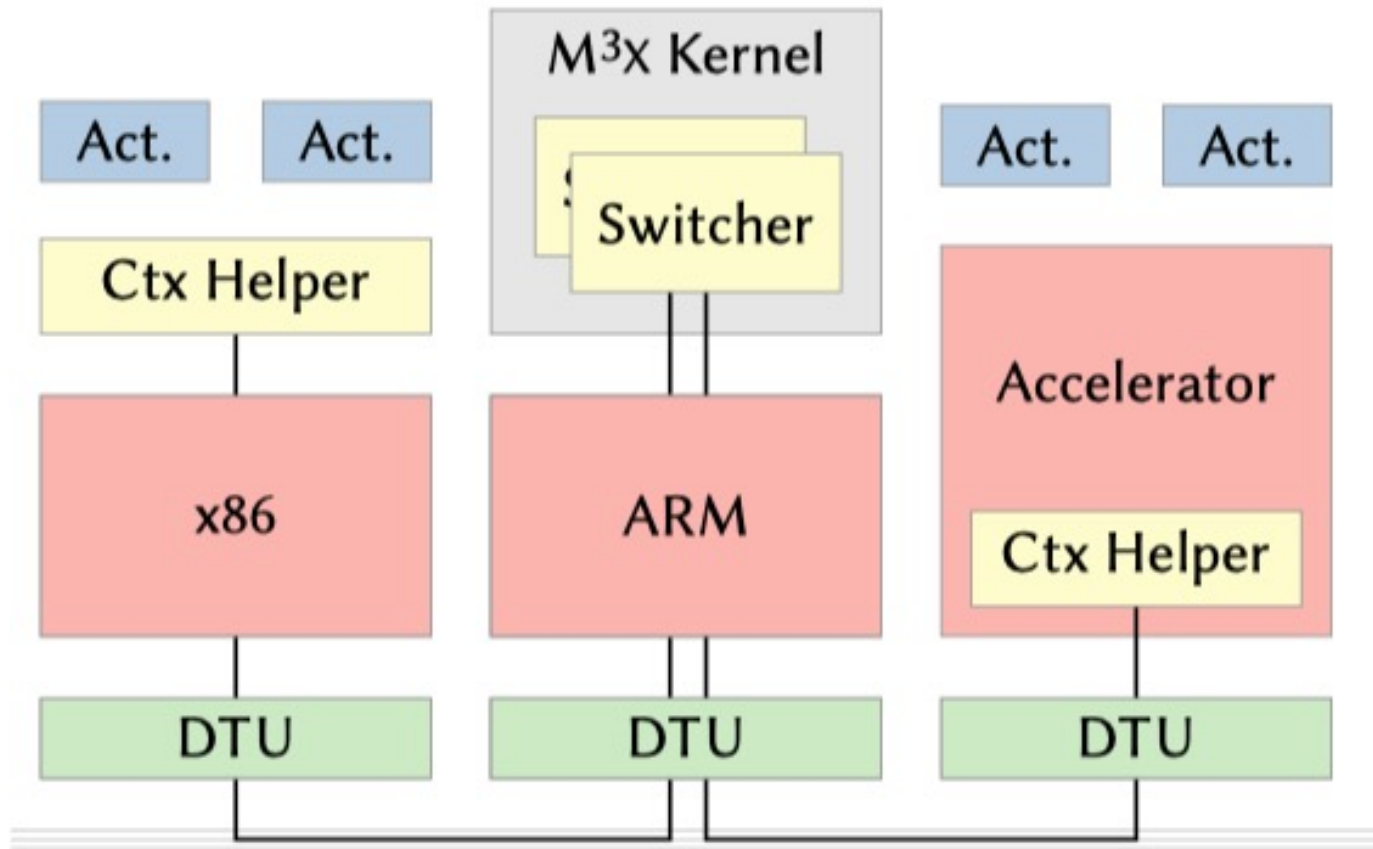
Accelerator Chains: Results (PCIe-like Latency)



Accelerator Chains: Results

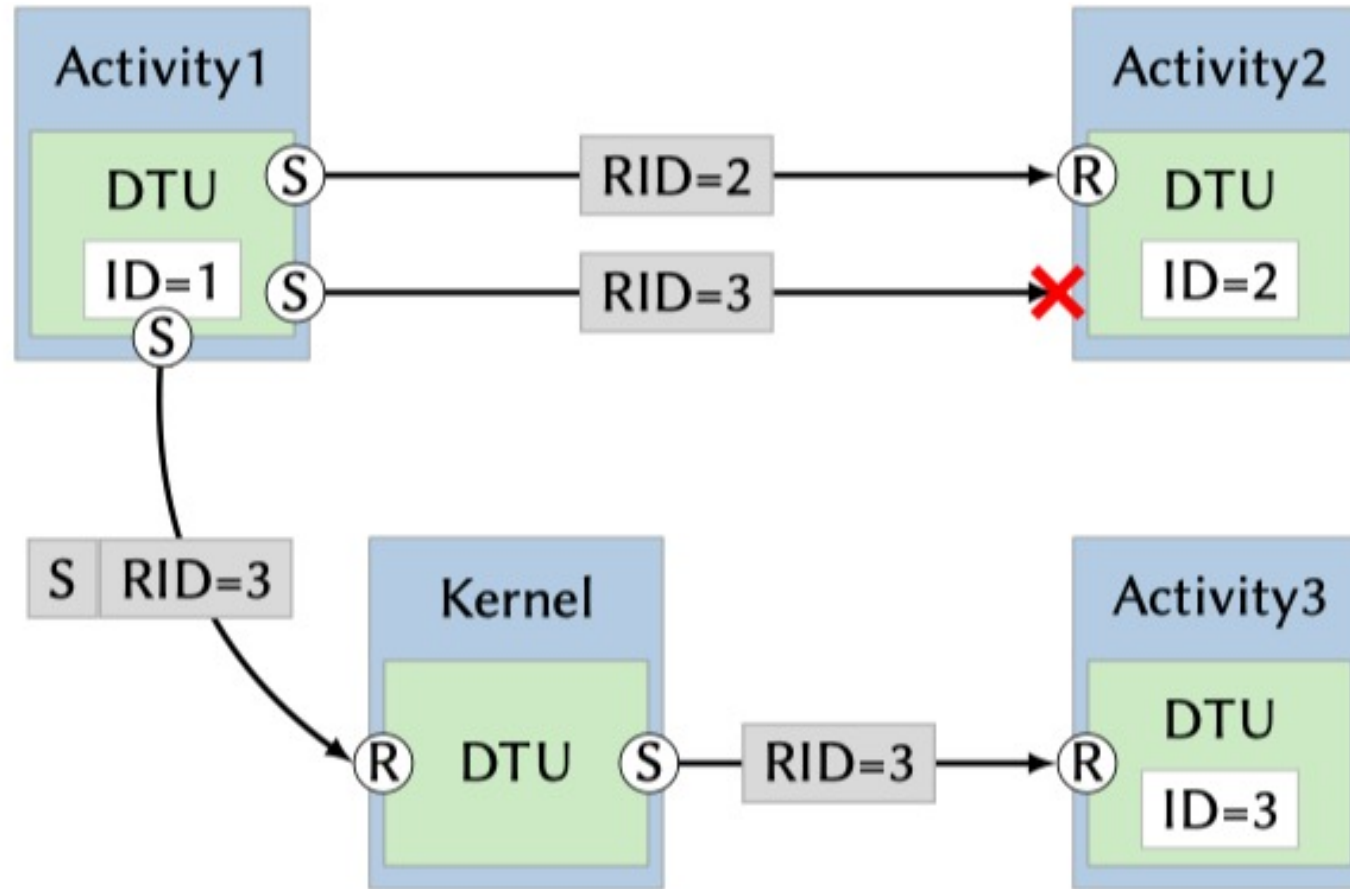


Fast-Path Communication vs. Context Switching



- Kernel handles complex part
 - Schedules/migrates activities
 - Initiates context switches
- Helper on user tiles implements save/restore
 - General purpose tiles: Software helper
 - Accelerator tiles: Helper implemented in hardware as part of ASM

Combining Fast-Path Communication with Context Switching



Conclusion

- M 3 X enables autonomous accelerators and combines fast-path communication with context switching:
 - Adding uniform interface to compute units
 - Using simple and generic protocols
 - Adding lightweight component to accelerators
- Reduces CPU load by a factor of 30
- Retains advantages of fast-path communication
- Uses hardware resource efficiently